

# AVL trees

# Dynamic set ADT

A **dynamic set ADT** is a structure that stores a set of elements. Each element has a (unique) **key** and **satellite data**. The structure supports the following operations.

**Search( $S, k$ )** Return the element whose key is  $k$ .

**Insert( $S, x$ )** Add  $x$  to  $S$ .

**Delete( $S, x$ )** Remove  $x$  from  $S$  (the operation receives a pointer to  $x$ ).

**Minimum( $S$ )** Return the element in  $S$  with smallest key.

**Maximum( $S$ )** Return the element in  $S$  with largest key.

**Successor( $S, x$ )** Return the element in  $S$  with smallest key that is larger than  $x.key$ .

**Predecessor( $S, x$ )** Return the element in  $S$  with largest key that is smaller than  $x.key$ .

# Motivation


- In a binary search tree, all operation take  $\Theta(h)$  time in the worst case, where  $h$  is the height of the tree.
- The optimal height of a binary search tree is  $\lfloor \log n \rfloor$ .
- Even if we start with a balanced tree, insertion/deletion operations can make the tree unbalanced.
- An **AVL tree** is a special kind of a binary search tree, which is always kept balanced.

# AVL tree

An **AVL tree** is a binary search tree such that for every node  $x$ ,

$$|\text{height}(x.\text{left}) - \text{height}(x.\text{right})| \leq 1$$

(we assume that  $\text{height}(\text{NULL}) = -1$ )



AVL trees are named after their inventors, Georgy Adelson-Velsky and Evgenii Landis.

# The height of an AVL tree

## Theorem

*The height of an AVL tree is  $\Theta(\log n)$ .*

- Let  $n_k$  be the minimum number of nodes in an AVL tree with height  $k$ .
- $n_0 = 1$ .




# The height of an AVL tree

## Theorem

*The height of an AVL tree is  $\Theta(\log n)$ .*

- Let  $n_k$  be the minimum number of nodes in an AVL tree with height  $k$ .
- $n_0 = 1$ .
- $n_1 = 2$ .




# The height of an AVL tree

## Theorem

*The height of an AVL tree is  $\Theta(\log n)$ .*

- Let  $n_k$  be the minimum number of nodes in an AVL tree with height  $k$ .
- $n_0 = 1$ .
- $n_1 = 2$ .
- $n_2 = 2 + 1 + 1 = 4$ .




# The height of an AVL tree

## Theorem

*The height of an AVL tree is  $\Theta(\log n)$ .*

- Let  $n_k$  be the minimum number of nodes in an AVL tree with height  $k$ .
- $n_0 = 1$ .
- $n_1 = 2$ .
- $n_2 = 2 + 1 + 1 = 4$ .
- $n_3 = 4 + 2 + 1 = 7$ .
- $n_k = n_{k-1} + n_{k-2} + 1$ .





# The height of an AVL tree

- $n_k = F_{k+3} - 1$  where  $F_k$  is the  $k$ -th Fibonacci number.

The proof is by induction:

Base:  $n_0 = 1 = F_3 - 1$ ,  $n_1 = 2 = F_4 - 1$ .

$$n_k = n_{k-1} + n_{k-2} + 1 = (F_{k+2} - 1) + (F_{k+1} - 1) + 1 = F_{k+3} - 1$$

- It is known that

$$F_k = \frac{\left(\frac{1+\sqrt{5}}{2}\right)^k - \left(\frac{1-\sqrt{5}}{2}\right)^k}{\sqrt{5}} \geq \frac{\left(\frac{1+\sqrt{5}}{2}\right)^k - 1}{\sqrt{5}}$$

- Therefore,

$$n_k \geq \frac{\left(\frac{1+\sqrt{5}}{2}\right)^{k+3} - 1}{\sqrt{5}} - 1 \geq \frac{\left(\frac{1+\sqrt{5}}{2}\right)^{k+3}}{\sqrt{5}} - 1.45$$

$$k \leq \log_{(1+\sqrt{5})/2} \left( \sqrt{5}(n_k + 1.45) \right) - 3 \leq 1.441 \log n_k.$$

- The height of an AVL tree with  $n$  nodes is  $\leq 1.441 \log n$ .

# Implementation

- A node in an AVL tree has the same fields defined for binary search tree (key, left, right, p).
- Additionally, every node has a field **height** which stores the height of the node.

# Searching



- The Search operation is handled exactly like in regular binary search trees.
- Time complexity:  $\Theta(h) = \Theta(\log n)$ .

# Insertion/Deletion

- Insertion and deletion are done by first applying the insertion/deletion algorithm of binary search trees.
- After the insertion/deletion, the tree may not be balanced, so we need to correct it.
- The time complexity of insertion/deletion in AVL tree is  $\Theta(\log n)$ .

# Insertion



After a new leaf is inserted, the height of some of its ancestors increase by 1. The heights of the other nodes are unchanged.



# Insertion



The height fields of the nodes can be updated by going up the tree from the inserted leaf, and for each ancestor  $v$  of the leaf perform

$$v.\text{height} = 1 + \max(v.\text{left}.\text{height}, v.\text{right}.\text{height})$$




# Insertion

Some of the ancestors of the leaf may become unbalanced.




- Let  $x$  be the lowest node on the path from the new leaf to the root which is unbalanced (if  $x$  doesn't exist we are done).
- There are 4 cases. In **Case 1** suppose that the new leaf is in the subtree of  $x.\text{left}.\text{left}$ .
- Let  $y = x.\text{left}$ .
- Let  $h$  be the height of the right child of  $x$ .
- Since  $x$  was balanced before the insertion, the height of  $y$  before the insertion was  $h - 1/h/h + 1$ .







- The insertion either increases the height of a node or doesn't change the height. Since  $x$  is now unbalanced, the only possible case is that the height of  $y$  is  $h + 1$  before the insertion, and  $h + 2$  afterward.
- The height of  $x$  before the insertion is  $h + 2$ , and  $h + 3$  afterward.





- The insertion either increases the height of a node or doesn't change the height. Since  $x$  is now unbalanced, the only possible case is that the height of  $y$  is  $h + 1$  before the insertion, and  $h + 2$  afterward.
- The height of  $x$  before the insertion is  $h + 2$ , and  $h + 3$  afterward.





- After the insertion,  $y$  has a child with height  $h + 1$ . This child must be the left child  $A$ . The height of  $A$  before the insertion is  $h$ .
- The height of the right child of  $y$  is  $h$ .





- To fix the imbalance of  $x$ , perform the following operation called **right rotation**.
- The rotation operation doesn't change the inorder of the nodes. Therefore, the new tree is a valid binary search tree.





- The heights of  $x$  and  $y$  after the rotations are  $h + 1$  and  $h + 2$ .
- After the rotation,  $x$  and  $y$  are balanced.





- Assume  $x$  had a parent the before the rotation, and denote it by  $u$ .
- Let  $w$  be the sibling of  $x$  before the rotation.
- The height of  $w$  is  $h + 1/h + 2/h + 3$ .
- If the height of  $w$  is  $h + 1$ , then  $u$  is unbalanced after the insertion (before the rotation).



- After the rotation, the height the sibling of  $w$  (node  $y$ ) is  $h + 2$ , which is equal to the height of the sibling of  $w$  before the rotation. Therefore,  $u$  is balanced.
- The height of  $u$  after the rotation is the same as the height before the insertion. Repeating these arguments, every ancestor of  $u$  is balanced and has same height as before the insertion.





# Example





In this example,  $h = 0$ .




- In **Case 2**, suppose that the new leaf is in the subtree of  $x.\text{left}.\text{right}$ .
- Let  $y = x.\text{left}$ .
- Performing a rotation on  $x$  and  $y$  does not work.



- Let  $z = y.\text{right}$ . Perform a **double rotation** on  $x, y, z$ .
- The double rotation doesn't change the inorder of the nodes.
- After the double rotation,  $x, y, z$  are balanced. Moreover, the height of  $z$  is the same as the height of  $x$  before the insertion, and therefore all ancestors of  $z$  are balanced.




- **Case 3** is when the new leaf is in the subtree of  $x.\text{right}.\text{right}$ , and **Case 4** is when the new leaf is in the subtree of  $x.\text{right}.\text{left}$ .
- Case 3 and Case 4 are symmetric to Case 1 and Case 2.
- Case 3 is shown below.




# Deletion


- After a node is deleted, the heights of some of its ancestors decrease by 1. The heights of the other nodes are unchanged.
- A single ancestor of the deleted node can become unbalanced.




- Let  $x$  be the unbalanced node (if  $x$  doesn't exist we are done).
- Assume that the deleted node is in the subtree of  $x.\text{right}$ .
- Let  $y = x.\text{left}$ .
- Since  $x$  was balanced before the deletion, the height of  $x.\text{right}$  before the deletion was  $h - 1/h/h + 1$ .




- The deletion either decreases the height of a node or doesn't change the height. Since  $x$  is now unbalanced, the only possible case is that the height of  $x.\text{right}$  is  $h - 1$  before the insertion, and  $h - 2$  afterward.
- The height of  $x$  is  $h + 1$  (the insertion doesn't change the height).



- The deletion either decreases the height of a node or doesn't change the height. Since  $x$  is now unbalanced, the only possible case is that the height of  $x.\text{right}$  is  $h - 1$  before the insertion, and  $h - 2$  afterward.
- The height of  $x$  is  $h + 1$  (the insertion doesn't change the height).




- Since  $y$  has height  $h$  and it is balanced, one of  $y$ 's children has height  $h - 1$  and the other child has height  $h - 1/h - 2$ .
- In **Case 1**, assume that the height of  $y.\text{left}$  is  $h - 1$ , and the height of  $y.\text{right}$  is  $h - 1/h - 2$ .







- In Case 1 we perform a right rotation on  $x$  and  $y$ .
- After the rotation, the height of  $x$  is  $h/h - 1$  and the height of  $y$  is  $h + 1/h$ .
- After the rotation,  $x$  and  $y$  are balanced.



- Assume  $x$  had a parent the before the rotation, and denote it by  $u$ .
- Let  $w$  be the sibling of  $x$  before the rotation.
- The height of  $w$  is  $h/h+1/h+3$ .
- $u$  is balanced after the insertion (before the rotation).



- After the rotation,  $u$  can become unbalanced. This occurs when the height of  $w$  is  $h + 2$ , and the height of  $y$  after the rotation is  $h$ .
- If the height of  $w$  is  $h$ , and the height after the rotation is  $h$ , then the height of  $u$  is  $h + 2$  and  $h + 1$  afterwards. This can cause an imbalance in an ancestor of  $u$ .
- In the two cases above, additional rotation is needed.



- In **Case 2**, assume that the height of  $y.\text{left}$  is  $h - 2$ , and the height of  $y.\text{right}$  is  $h - 1$ .
- Let  $z = y.\text{right}$ .
- In this case we perform a double rotation of  $x, y, z$ .

